

DeliveryBrief AI Collaboration Note

Work delegated to AI, corrections, verification and decisions I owned

AI COLLABORATION NOTE

Prepared by Elvis

7 September 2026

My responsibility

I used AI to accelerate analysis, scaffolding and review. I remain responsible for the workflow choice, source permissions, code, evaluation, claims and submission. I will not present an AI suggestion as an observed result.

7 September Quest interpretation and scope

Tool and model. Codex coding agent.

Delegated work. I asked the agent to extract the Quest requirements, compare project directions and translate the scoring rubric into a build strategy.

Accepted. I retained the recommendation to choose one recurring workflow with a measurable baseline, two real integrations, human approval and explicit failure cases.

Rejected or corrected. I did not treat a polished demo as proof of operational value. I also replaced the original OpenAI recommendation after confirming that I already had Anthropic API credit.

Verification. I checked the recommendations against the supplied Quest and role description.

Decision I owned. I selected the workflow, target user, sources, public demo and funded model provider.

7 September system implementation

Tool and model. Codex coding agent.

Delegated work. I asked the agent to scaffold the Python application, schemas, read-only source adapters, model integration, validation, storage, exports, cases, tests and document drafts.

Accepted. I retained the separation between a credential-free sample and live integrations. Every factual item carries evidence IDs, and external sending remains outside the system.

Rejected or corrected. No field metric, adoption claim or user quote was generated. The documents identify the evidence still required. During browser testing, I reviewed the sample output and corrected duplicate topic handling and a phone-pattern false positive.

Verification. Twelve automated tests, Ruff, mypy, the deterministic case runner, the Streamlit health endpoint and the complete browser workflow were executed.

Decision I owned. I will provide credentials, approve the anonymized evidence, conduct the manager session and review every final statement before submission.

7 September evaluation and cost control

Tool and model. Codex coding agent and Claude Haiku 4.5.

Delegated work. I asked the agent to run capped model evaluations, compare the direct-prompt baseline with the DeliveryBrief workflow and stop Sonnet unless I explicitly approved more spend.

Accepted. The direct-prompt Haiku baseline passed 1 of 9 scored cases. The DeliveryBrief Haiku workflow passed 9 of 9 scored report cases with a \$0.30 cap and an actual estimated cost of \$0.025316.

Rejected or corrected. I did not accept the first 8-of-9 result as final. The failed cases led to fixes for strict Anthropic schemas, cross-repository evidence coverage and missing action-like evidence.

Verification. Focused regressions passed, the final full Haiku run passed, and tests, Ruff, mypy and the secret scan passed.

Artifacts

- Typed application and integration code
- Ten-case evaluation manifest and capped Haiku result
- Direct-prompt baseline result
- Public app-run exports and screenshot
- Automated tests and static-analysis configuration
- Decision record, runbook and submission-document sources

Continuation format

For every later AI-assisted task I will add the date, tool and model, delegated work, accepted output, rejected or corrected output, verification, decision I owned and artifact reference. Corrections will link to a commit, test or result file where possible.

Entries still required

- Live credential and integration verification
- Changes made after any external manager review
- Final document and demo-video review

Disclosure principle

The purpose of this note is to make the collaboration inspectable. I will describe AI's contribution accurately and avoid language that implies I completed work I cannot explain or verify.