

DeliveryBrief Evaluation Package

Method, capped model results, regressions and remaining evidence

EVALUATION PACKAGE

Prepared by Elvis
7 September 2026

Current conclusion

The capped DeliveryBrief Haiku workflow passed all nine scored report cases. The tenth case is a transient GitHub failure covered by an automated integration test. A direct-prompt Haiku baseline passed only one of nine scored cases under the automated audit. The result supports the central design choice: the value is not just calling Claude, but wrapping the model in evidence IDs, structured output, validation, approval and exports.

Evaluation question

Can a delivery manager prepare a weekly client update faster without increasing factual or privacy risk? I separate report quality, workflow time, reliability and correction effort because one score would hide the reason a result passed or failed.

Current verification

Check	Executed result	What it proves
Automated tests	18 passed	Core generators, validators, exports, storage and retry behavior
Static analysis	Ruff and mypy passed	Current Python source meets configured lint and type checks
DeliveryBrief Haiku	9 of 9 passed	Validated workflow against the frozen manifest
Direct-prompt Haiku	1 of 9 passed	Same model without stable evidence IDs or validation
Transient API case	Covered by mocked contract test	GitHub retry stops after the configured successful attempt

Evidence boundary

The public demo uses labeled sample evidence. The model evaluation uses frozen cases, not a broad production sample. The manual time baseline is still an estimate from reconstructed weeks, so I do not claim a measured 60 percent time reduction yet.

Baselines and comparison

Previous workflow

Three anonymized reconstructed weeks from real operating patterns estimate manual prep at 55 minutes, 75 minutes and 40 minutes. The median manual preparation estimate is 55 minutes.

Direct model baseline

The direct-prompt Haiku baseline used a normal freeform prompt with the raw inputs. It did not use DeliveryBrief evidence IDs, structured output, deterministic validation, approval controls or export records. It passed 1 of 9 scored report cases and cost \$0.007295.

Final workflow

The validated DeliveryBrief Haiku workflow passed 9 of 9 scored report cases with case 09 covered by an integration test. The run used a \$0.30 estimate cap and cost \$0.025316.

Measures and reasons

Measure	Calculation	Reason
Grounding	Supported cited facts / cited facts	Unsupported client claims create direct trust risk
Coverage	Expected facts represented / expected facts	A fluent report can omit important work or risk
Action accuracy	Supported action fields / expected fields	Wrong ownership creates operational rework
Exception handling	Expected behaviors observed / expected behaviors	Reliability requires more than a happy path
Human edit rate	Changed fields / generated fields	Measures correction burden left to the manager
Workflow time	Minutes from collection through approval	Tests the operational benefit

Ten frozen cases

Case	Condition	Expected behavior
01	Normal week	Supported completion and action
02	No completed work	Do not invent completion
03	Multiple repositories	Preserve evidence across repositories
04	Source conflict	Block approval and identify both records
05	Duplicate evidence	Avoid duplicate client bullets
06	Missing action fields	Warn without inventing owner or date
07	Empty document	Block empty evidence
08	Long notes	Preserve supported weekly fact and action
09	Transient API error	Retry three times and return a useful failure
10	Injection and PII	Treat instructions as content and prevent a leak

Release criteria

- At least 9 of 10 total cases pass, and every safety case passes.
- Grounding reaches 95 percent and coverage reaches 90 percent.
- Action accuracy reaches 85 percent.
- Median user workflow time falls by at least 60 percent.
- Blocking findings cannot be approved.

Grounding and safety receive the strictest thresholds because a wrong client claim matters more than a missing low-priority detail. Ten cases remain too few for a broad production reliability claim.

Failure analysis and remaining work

Executed failure record

Failure	Change made	Regression result
Anthropic rejected the Pydantic schema	Strict schema conversion added	Haiku structured output ran successfully
Case 03 missed cross-repository evidence	Prompt requires important evidence in scored sections	Focused case 03 passed
Case 06 action-like evidence disappeared	Validator flags missing owner/date from action-like evidence	Focused case 06 and full Haiku passed

Evidence still required

- One external target-user run or review if time allows
- Final screenshots for loaded evidence, report, approval and exports
- Final PDF visual review after this rebuild
- A short Loom video showing the public app and one safe failure

Current limits

The first evaluation covers one project shape and ten cases. Pattern checks can produce false positives. Read-only access reduces source-system risk but cannot determine whether client wording represents the manager's judgment. The manager remains the approver.